

**Operating Systems:**  
**Chaps 2, 3, 4**  
**Scheduling Algorithms**  
**CS400**

Week 8: 3<sup>rd</sup> March

Spring 2020

Oliver BONHAM-CARTER

# What's Been Covered

- We have learned about:
  - Processes
    - Creating
    - Destroying
    - Blocking
  - Threads
    - Creating
    - Destroying
  - Semaphores
  - Memory management institutions
    - Maps
    - Paging systems



# *Bursts of Usage*

- A “burst”: A brief stretch of “run as fast as you can go until you cannot continue for some reason.
- A CPU bursts when it is executing instructions;
  - An I/O system bursts when it services requests to fetch information.
  - Each component operates until it cannot continue (for some reason).

# *Bursts of CPU and IO Usage*

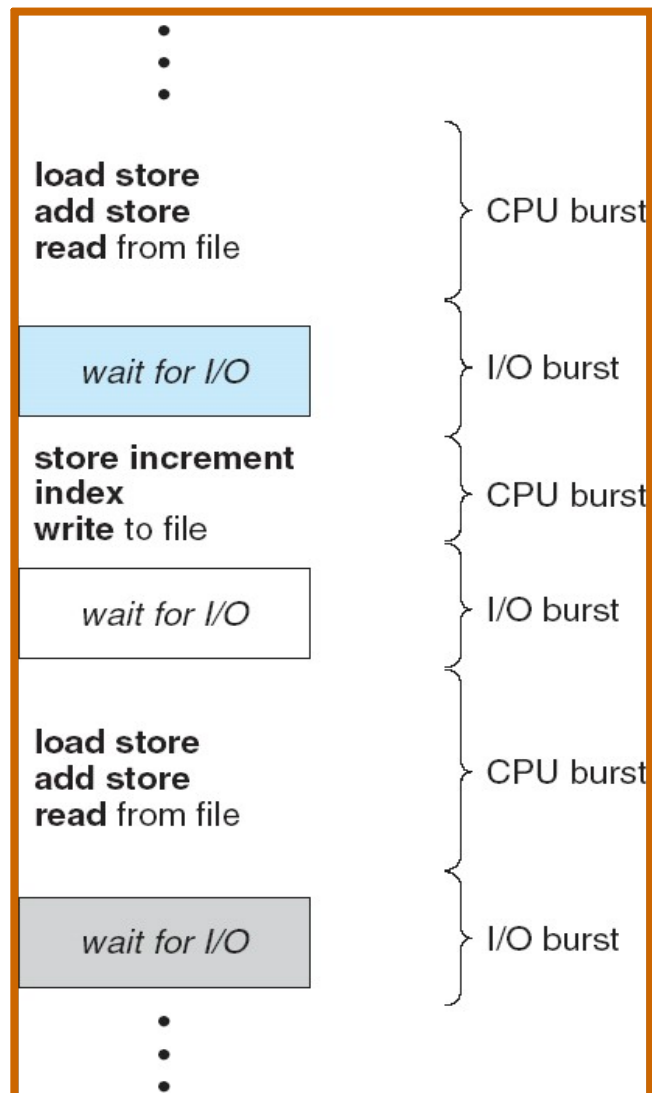
- Start of CPU burst
  - A CPU can run instructions from cache until it needs to fetch more instructions or data from memory.
- Ends of CPU burst
- Starts the I/O burst.
  - The I/O burst read or writes data until the requested data is read/written or the space to store it cache runs out.
- End of an I/O burst.

# Controlling *Bursts*: The Magic of the OS

- *“The magic of an OS is the act of managing and scheduling these activities to maximize the use of the resources and minimize the wait and idle times for processes.”*




# Alternating Sequence: CPU and I/O *Bursts*



- A typical execution of a program is made-up by bursts
- CPU: runs only one thing at a time
- I/O: Reads or writes to complete the instruction
- *All or nothing type of system*

# Histogram of CPU-*Burst* Times



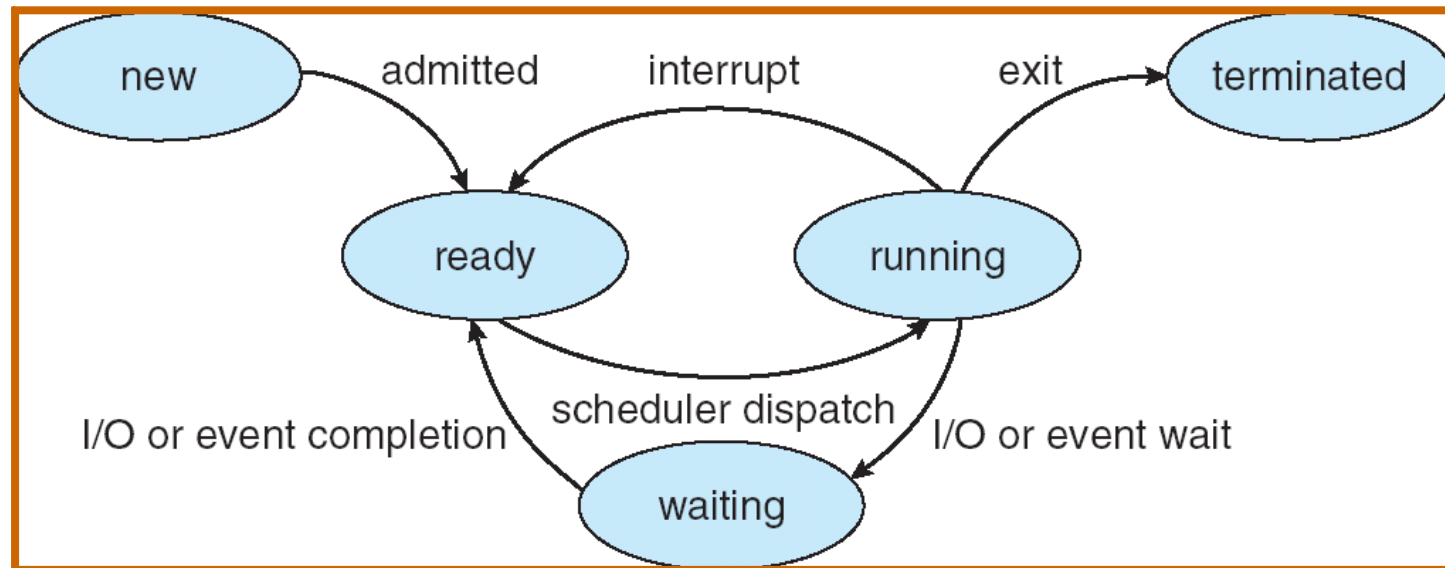
CPU bursts tend to have a frequency curve similar to the exponential curve shown above. It is characterized by a large number of short CPU bursts and a small number of long CPU bursts. An I/O-bound program typically has many short CPU bursts; a CPU-bound program might have a few long CPU bursts.

# States of a Process

- **READY** - The process is waiting to be assigned to a processor.
- **RUNNING** - Instructions are being executed.
- **WAITING** - The process is waiting for some event to occur (such as an I/O completion or reception of a signal).
- **TERMINATED** - The process has finished execution.



# CPU Scheduler: Handling *Bursts*



- The CPU scheduler selects from among the processes in memory that are ready to execute and allocates the CPU to one of them.

# CPU Scheduler

- CPU scheduling is affected by the following set of circumstances:
  - A process switches from **running** to **waiting** state
  - A process switches from **running** to **ready** state
  - A process switches from **waiting** to **ready** state
  - A processes switches from **running** to **terminated** state
- Circumstances 1 and 4 are non-preemptive; they offer **no scheduling choice**
- Circumstances 2 and 3 are pre-emptive; they **can be scheduled**

# Dispatcher

- The dispatcher module gives control of the CPU to the process selected by the short-term scheduler; involving:
  - switching context
  - switching to user mode
  - jumping to the proper location in the user program to restart that program
- The dispatcher needs to run as fast as possible, since it is invoked during process context switch
- The time it takes for the dispatcher to stop one process and start another process is called *dispatch latency*

# Scheduling Criteria: Algorithms

How many shortest-length paths are there to get from your house to the doughnut shop?

4 up's  
7 right's

$\binom{11}{7} = \binom{11}{4} = 330$  paths

$\binom{n}{k} = \frac{n!}{(n-k)!k!}$

$e^{i\pi} + 1 = 0$

$B_4$

Onto

One-to-One

There are six dogs to give 13 tacos. Use a 'stars and bars' diagram to illustrate the first and sixth dog get 3 tacos, the second dog gets none, the third dog gets 5 and the fourth dog gets one.

☆☆☆||☆☆☆☆☆|☆||☆☆☆

$A = \{2, 4, 10, \frac{1}{2}\}$

$(A \cup B \cup C) \cup (A \cap B \cap C)$

$P, Q, R$

P	Q	R	P ∨ Q	P ∨ R	(P ∨ Q) ∧ (P ∨ R)
T	T	T	T	T	T
T	T	F	T	T	T
T	F	T	T	T	T
T	F	F	T	F	F
F	T	T	T	T	T
F	T	F	T	F	F
F	F	T	F	T	F
F	F	F	F	F	F

Find  $7 + 12 + 17 + 22 + \dots + 342$

$S_n = 7 + 12 + 17 + 22 + \dots + 342$

$2S_n = 342 + 337 + 332 + 327 + \dots + 7$

$2S_n = 349 + 68$

$S_n = \frac{349 + 68}{2}$

$S_n = 11866$

Original:  
 $\exists x \forall y (x \geq 2y \rightarrow x > y + 1)$

Converse:  
 $\exists x \forall y (x > y + 1 \rightarrow x \geq 2y)$

Negation:  
 $\neg [\exists x \forall y (\neg (x \geq 2y) \vee x > y + 1)]$

Contrapositive:  
 $\exists x \forall y (x \leq y + 1 \rightarrow x < 2y)$

$v = e + f = 2$

P.I.E. Example:

$6! - \left[ \binom{6}{1}5! - \binom{6}{2}4! + \binom{6}{3}3! - \binom{6}{4}2! + \binom{6}{5}1! \right]$

- Different CPU **scheduling algorithms** have different properties
- The choice of a particular algorithm may favor one class of processes over another
- In choosing which algorithm to use, the properties of the various algorithms should be considered


# Scheduling Criteria:

## *response, waiting and turnaround*

- Criteria for comparing CPU scheduling algorithms may include the following
  - **CPU utilization** – percent of time that the CPU is busy executing a process
  - **Throughput** – number of processes that are completed per time unit
  - **Response time** – amount of time it takes from when a request was submitted until the first response occurs (but not the time it takes to output the entire response)
  - **Waiting time** – the amount of time before a process starts after first entering the ready queue (or the sum of the amount of time a process has spent waiting in the ready queue)
  - **Turnaround time** – amount of time to execute a particular process from the time of submission through the time of completion

# Optimization Criteria

- It is desirable to
  - Maximize CPU utilization
  - Maximize throughput
  - Minimize turnaround time
  - Minimize start time
  - Minimize waiting time
  - Minimize response time
- In most cases, we strive to optimize the average measure of each metric
- In other cases, it is more important to optimize the minimum or maximum values rather than the average




# What Makes Scheduling Happen?

- Algorithms (*programming recipes for specific tasks*)
- Scheduling Algorithms:
  - First Come, First Served (FCFS)
  - Shortest Job First (SJF)
  - Priority
  - Round Robin (RR)

# Algorithm: First Come, First-Served (1)

- Process: Burst Time
  - P1: 24
  - P2: 3
  - P3: 3
- With FCFS, the process that requests the CPU first is allocated the CPU first
- Case #1: Suppose that the processes arrive in the order: P1 , P2 , P3

The Gantt Chart for the schedule is:




- **Waiting time** for P1 = 0; P2 = 24; P3 = 27
- **Average waiting time:**  $(0 + 24 + 27)/3 = 17$
- **Average turn-around time:**  $(24 + 27 + 30)/3 = 27$



# Algorithm: **First Come, First-Served** (2)

- Case #2: Suppose that the processes arrive in the order: P<sub>2</sub> , P<sub>3</sub> , P<sub>1</sub>
- The Gantt chart for the schedule is:



- **Waiting time** for P<sub>1</sub> = 6; P<sub>2</sub> = 0; P<sub>3</sub> = 3
- **Average waiting time:**  $(6 + 0 + 3)/3 = 3$  (Much better than Case #1: 17)
- **Average turn-around time:**  $(3 + 6 + 30)/3 = 13$  (case #1: 27)
- Case #1 is an example of the **convoy effect**;
  - all the other processes wait for one long-running process to finish using the CPU
- This problem results in lower CPU and device utilization; Case #2 shows that higher utilization might be possible if the short processes were allowed to run first
- The FCFS scheduling algorithm is non-preemptive (non-scheduling)
- Once the CPU has been allocated to a process, that process keeps the CPU until it releases it either by terminating or by requesting I/O
- It is a troublesome algorithm for time-sharing systems (*users and devices are hard to control!*)


# Algorithm: Shortest Job First

- The SJF algorithm associates with each process the length of its next CPU burst
- When the CPU becomes available, it is assigned to the process that has the smallest next CPU burst (in the case of matching bursts, FCFS is used)
- Two schemes:
  - **Nonpreemptive** – once the CPU is given to the process, it cannot be preempted until it completes its CPU burst
  - **Preemptive** – if a new process arrives with a CPU burst length less than the remaining time of the current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (SRTF)

## Algorithm: **Shortest Job First** Ex1 (simultaneous arrival)

Process	Arrival Time	Burst Time
$P_1$	0.0	6
$P_2$	0.0	4
$P_3$	0.0	1
$P_4$	0.0	5

- SJF (non-preemptive, simultaneous arrival)




- **Average waiting time** =  $(0 + 1 + 5 + 10)/4 = 4$
- **Average turn-around time** =  $(1 + 5 + 10 + 16)/4 = 8$

## Algorithm: SJF Ex2 (non-preemptive, varied arrival times)

- Non-preemptive (no stalling)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

- SJF (non-preemptive, varied arrival times)



- Average waiting time**

$$= (P_1 + P_2 + P_3 + P_4) / 4$$

$$= ((0 - 0) + (8 - 2) + (7 - 4) + (12 - 5)) / 4$$

$$= (0 + 6 + 3 + 7) / 4 = 4$$

- Average turn-around time:**


$$= ((7 - 0) + (12 - 2) + (8 - 4) + (16 - 5)) / 4$$

$$= (7 + 10 + 4 + 11) / 4 = 8$$

# Preemptive SJF (shortest-remaining time first)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

- SJF (preemptive, varied arrival times)



- **Average waiting time**

$$\begin{aligned}
 &= ( [(0 - 0) + (11 - 2)] + [(2 - 2) + (5 - 4)] + (4 - 4) + (7 - 5) ) / 4 \\
 &= 9 + 1 + 0 + 2 / 4 \\
 &= 3
 \end{aligned}$$

- **Average turn-around time** =  $(16 + 7 + 5 + 11) / 4 = 9.75$

# Algorithm: **Priority** (1)

- The SJF algorithm is a *special case* of the general priority scheduling algorithm
- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer = highest priority)
- Priority scheduling can be either *preemptive* or *non-preemptive*:
  - A **preemptive** approach will preempt the CPU if the priority of the newly-arrived process is higher than the priority of the currently running process
  - A **non-preemptive** approach will simply put the new process (with the highest priority) at the head of the ready queue

## Algorithm: **Priority** (2)

- SJF is a priority scheduling algorithm where priority is the predicted next CPU burst time
- The main problem with priority scheduling is *starvation*, that is, low priority processes may never execute
- A solution is *aging*; as time progresses, the priority of a process in the ready queue is increased

# Algorithm: Round Robin


- In the round robin algorithm, each process gets a small unit of CPU time (a time quantum), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are  $n$  processes in the ready queue and the *time quantum* is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once. *No process waits more than  $(n-1)q$  time units.*
- Performance of the round robin algorithm
  - $q$  large --> Simply use: the First Come, First Served algorithm
  - $q$  small -->  $q$  must be greater than the context switch time; otherwise, the overhead is too high
- Typically: 80% of the CPU bursts should be shorter than the time quantum
- *Def of time quantum: The period of time for which a process is allowed to run uninterrupted in a pre-emptive multitasking operating system*



## Round Robin Ex, Time Quantum = 20

Process	Burst Time
$P_1$	53
$P_2$	17
$P_3$	68
$P_4$	24

- The Gantt chart is:



0    20    37    57    77    97    117    121    134    154    162    <---- time

- Typically, higher average turnaround than SJF, but better response time
- Average waiting time**  

$$= ( [(0 - 0) + (77 - 20) + (121 - 97)] + (20 - 0) + [(37 - 0) + (97 - 57) + (134 - 117)] + [(57 - 0) + (117 - 77)] ) / 4$$

$$= (0 + 57 + 24) + 20 + (37 + 40 + 17) + (57 + 40) / 4$$

$$= (81 + 20 + 94 + 97) / 4$$

$$= 292 / 4 = 73$$
- Average turn-around time**  

$$= (134 + 37 + 162 + 121) / 4 = 113.5$$







# Nice Scheduling Demo

- <https://ess.cs.tu-dortmund.de/Software/AnimOS/CPU-Scheduling/>

▼ Definition

Operate

Help

Fn	Process name	Arrival	CPU burst	IO burst
 	Process1	5	4	9
 	Process2	3	4-7	6
 	Process3	1	7	4-5
<div> <div>+ Add</div> <div>Process name</div> <div>Arrival</div> <div>CPU burst</div> <div>IO burst</div> </div>				